

Himanshu Rane

✉ himanshu.rane005@gmail.com

☎ +49 15510826125

📍 Bahnhofstraße 81, Neu-Ulm

🌐 linkedin.com/in/himanshu-rane-

🐙 github.com/19himanshurane

🔗 himanshurane.com



📄 SUMMARY

Master's student in AI & Data Analytics at HNU Neu-Ulm with 6 months of hands-on experience as a Business and Data Analyst. I build production-thinking AI systems, a multi-agent critic pipeline, an LLM regression test suite and a RAG-based chatbot that cut processing time by 83%. Looking for a working student or internship role in AI Engineering or Data analytics.

🎓 EDUCATION

Master of Science – Artificial Intelligence and Data Analytics

Hochschule Neu-Ulm (HNU)

10/2024 – Present

Germany

Bachelor of Engineering – Electronics and Telecommunication Engineering

Jayawantrao Sawant College of Engineering

08/2020 – 06/2024

Pune, India

📁 PROJECTS

Quorum - Multi Agent LLM Verification System

- Designed a multi-agent system where three independently modeled LLM critics (accuracy, logic, completeness), each on a different model family, evaluate an output in parallel.
- Built a LangGraph-orchestrated adjudicator that resolves critic disagreements into a single verdict with a confidence score and supporting evidence.
- Diagnosed and fixed a thread-safety deadlock in concurrent provider client construction using a per-provider lock and cache.
- Shipped a Streamlit “Verdict Explorer” UI, a documented FastAPI backend, a full SQLite audit trail, one-command Docker Compose deployment, and a pytest suite covering every routing path.

EvalGate - LLM Evaluation & MLOps CI/CD Pipeline

- Built a CI/CD-style regression pipeline that tests LLM-powered features against a 50-case hand-verified golden dataset on every prompt or model change.
- Implemented multi-dimensional scoring (category accuracy, LLM-as-judge quality rating, latency, token cost) with statistical drift thresholds and a 7-run rolling average.
- Built a GitHub Actions workflow that posts pass/fail summaries and blocks merges on critical regressions, with Slack alerts.
- Took the golden-dataset pass rate from 84% to 100% across four real prompt iterations; a live test PR caught two CI-only bugs invisible in local testing.

LLM Cost Autopilot

- Built a scikit-learn complexity classifier that scores prompts into 3 tiers using lightweight text features, at 86% accuracy on held-out data.
- Router maps each tier to a model via a hot-reloadable config — routing changes take effect through the API with no redeploy.
- Async verifier re-runs requests through a reference model in the background, comparing answers and auto-escalating on disagreement.
- Escalated mismatches feed back into the training set, closing the loop for the next classifier retrain.
- Shipped a live Streamlit cost dashboard (actual vs. baseline spend, routing distribution, escalation rate) reading live from a FastAPI backend, containerized with Docker and covered by CI.

Meter Mate - AI Research Agent (RAG)

- Led a 5-person cross-functional team building an AI research agent with LangChain, ChromaDB, and LLaMA 3 to retrieve and structure information from large operational document corpora.
- Cut manual research time per query from 5 minutes to 30 seconds, an 83% reduction.
- Shipped a working Python + Streamlit prototype, iterated on stakeholder feedback, and delivered a production-ready system with documentation.
- Presented findings and the prototype directly to management.

Leaf Disease Detection

Research Author

- Co-authored a peer-reviewed paper evaluating AlexNet and GoogLeNet CNN architectures for potato leaf disease classification (healthy / early blight / late blight), reaching 86.35% accuracy on a withheld test set of ~3,000 segmented leaf images.
- Built the image preprocessing and segmentation pipeline (color/lightness/saturation masking in Lab and HSB space) and ran train/test split experiments across five ratios to validate results against overfitting.

 PROFESSIONAL EXPERIENCE

Business Analyst Intern

02/2024 – 07/2024

BitFlyer Technologies

Pune, India

- Built automated data collection and processing pipelines with Python and REST APIs, eliminating 8+ hours of manual work per week.
- Analyzed 47,000+ records to surface patterns and actionable insights for management decision-making.
- Developed Power BI dashboards enabling Sales and Operations to track KPIs independently, cutting the reporting cycle from 5 hours to 2 hours (60% faster).
- Contributed analysis and recommendations that shaped quarterly sales strategy.

 SKILLS

Tools & Platforms

GitHub, Claude Code, Cursor, Docker, n8n, Power BI, VSCode, Jira, SAP Signavio, Power Automate, Microsoft Office Suite,

Programming :

Python (3+ years), SQL, REST APIs, FastAPI

Data Analytics

DAX, Excel, Power BI, ETL Pipelines, Data Modelling, Data Quality

AI/ML & LLM Engineering

LLMs, LangChain, LangGraph, RAG, Prompt Engineering, AI Agents, LLM Evaluation, ChromaDB, Transformers, Attention mechanism, Deep Learning, MCP Servers, Machine Learning

 LANGUAGES

English

C1

German

B1

 CERTIFICATES

- IBM Python for Data Science, AI and Development, Coursera, 2025
- IBM Data Visualization and Dashboards with Excel and Cognos, Coursera, 2025
- Excel Essentials for Data Analytics, Coursera, 2025
- Deloitte Australia Data Analytics Job Simulation, Forage, 2025
- Goethe-Zertifikat A1 Start Deutsch 1, Goethe-Institut, 2025